

What Agentic AI taught us.

Part 3

**How to earn users' trust
in autonomous systems**



IXOR

digital impact created by talent

Why trust matters

Agentic AI can be smart, scalable, and autonomous, but it will fail without trust. No matter how capable a system is, users won't adopt it if they don't feel comfortable handing over control. Trust isn't just a bonus; it's the foundation for real-world adoption.

As agents become more independent, the challenge changes: it's no longer just about proving that a system works, but showing it can be trusted. Without trust, even the most advanced AI becomes shelfware.

Trust is what turns a working prototype into a daily tool. It's what makes people delegate tasks, rely on outcomes, and keep using the system - especially when they don't fully understand how it works.

This Impact Paper explores how to build that trust through transparency, predictability, and respectful design.



What we mean by trust

In Agentic AI, trust is not only about the agent itself, but also about the ecosystem around it.

Users feel safe when the agent provides valid answers, not just socially acceptable ones designed to please, but responses that are genuinely useful and reliable.

Trust also grows when the interaction feels natural and the agent clearly fits the business context.

Equally important is what happens behind the scenes: the team and the architecture. Knowing that experts have designed, host, and monitor the system makes trust sustainable.



Why trust is harder in Agentic AI

Agentic systems don't just respond to input: they plan, take action, use tools, and adapt. That flexibility is powerful, but also introduces complexity.

Key challenges include:

- **Unpredictability:** Small context changes can lead to very different outcomes.
- **Hidden decisions:** Memory, tools, and intermediate steps happen behind the scenes.
- **Lack of control:** Users may not know how to intervene or redirect the agent.
- **Shifting behaviour:** Learning systems may evolve in ways that confuse or unsettle users.

Classic software earns trust through reliability. Agentic AI must also manage initiative and uncertainty, making the trust challenge more dynamic and nuanced.

How to build trust into Agentic AI

Through our experience building production-ready agents, and insights from the experience, we recommend five core principles:

1. Be transparent, but relevant

Don't flood users with logs or token counts. Instead, offer timely, useful cues - like what the agent plans to do next, or why it picked a certain path.

2. Design for predictability

Use consistent language, outputs, and flows. Predictable agents feel smarter because they create less friction. Surprises break confidence.

3. Respect boundaries

Don't assume consent. Confirm before high-impact actions. Let users pause, reverse, or review what's happening. Respect in design builds respect in use.

4. Show confidence levels

Let users know when the agent isn't sure. Don't fake confidence.

5. Give users control points

Autonomy doesn't mean loss of control. Let users set boundaries, review actions, or intervene. Giving people ways to steer makes the system feel safer, even if they rarely use them.





What we learned about trust in agentic AI

From our own projects we've seen what really shapes trust in agentic systems. These lessons didn't come from theory, but from building and deploying agents in practice.

Honest answers build confidence

Trust starts with what the agent says. People value agents that don't fall back on socially desirable answers but give honest, valid help they can use. Reliability beats politeness.

Agents can earn trust

We noticed that trust was there when agents behaved consistently and delivered useful results. It wasn't automatic. Trust grew with every valid answer and every helpful action.

Humanity can trigger trust, but is hard to achieve

A human touch makes a big difference. Tone, empathy and natural interaction styles make systems approachable. But this is also where the most work is needed. Getting "human" right is harder than it looks.

The team behind the agent matters

Trust is not just in the interface. It comes from the people who design, train and host the system. Expertise and support structures give users confidence that the system won't fail them.

Architecture matters. A fragile setup breaks trust quickly. But when systems are secure, monitored and well designed, users feel safe to rely on them.

Business context is key

Agents must understand the business they serve. When they fit into existing processes, use the right language and align with real goals, adoption is faster.

Trust in agentic AI never comes from one factor alone. It's the mix of valid answers, human-centred design, expert teams, solid architecture and business alignment that makes people confident enough to rely on agents every day.

Conclusion: trust enables autonomy

Agentic AI won't win adoption because it's smart. It will win because people feel safe using it. That means designing systems that communicate clearly, act predictably, and know when to pause, ask, or admit doubt. Trust isn't a bonus feature. It's the intersection between autonomy and usability.

What can Ixor do?

At Ixor, we have already brought several Agentic AI projects into production, and are currently working on more.

This gives us a strong foundation of real-world experience, not just theoretical knowledge.

Our teams understand what works, where the pitfalls are, and how to translate the promise of Agentic AI into practical, scalable solutions.

Whether you're exploring the potential or ready to start building, we can help you move forward with clarity and confidence.

**Want to explore this in your organisation?
Book an Agentic AI Day. It's a practical session to turn AI ideas into action.**

More info here: ixor.be/agentic-ai-days



For any questions about Agentic AI, contact **Jan Pieter Cootjans**, our in-house expert.

janpieter.cootjans@ixor.be
+32 (0)475 20 57 67
www.ixor.be

Not every process needs an agent.

In our next Impact Paper, we'll help you spot where Agentic AI adds value - and where it doesn't. We'll share a practical checklist to help you identify the right use cases and assess the potential value for your organisation.

IXOR

digital impact created by talent